

Akshya Dhiman

Greater Noida | akshya1108dhiman@gmail.com | +91-7037843589 | [LinkedIn](#) | [GitHub](#)

SUMMARY

Entry-level Machine Learning Engineer with hands-on experience building end-to-end AI systems, including RAG pipelines, recommendation engines, and ETL workflows. Proficient in Python, SQL, and modern ML/LLM frameworks (LangChain, FAISS, PyTorch, Scikit-learn), with a focus on grounded LLM applications, feature engineering, and scalable inference.

TECHNICAL SKILLS

Programming: Python, SQL

Machine Learning: Regression, Classification, Feature Engineering, Model Evaluation (Precision, Recall, F1), Hyperparameter Tuning, Scikit-learn

Deep Learning: PyTorch, CNN, RNN

AI / LLM Systems: LangChain, Retrieval-Augmented Generation (RAG), FAISS, Ollama, Local LLMs, Prompt Engineering

Data Engineering: ETL Pipelines, Data Modeling, Data Warehousing (Snowflake)

Data Analysis: Pandas, NumPy, Exploratory Data Analysis (EDA), Data Cleaning

Databases: MySQL, Snowflake

Visualization & BI: Matplotlib, Seaborn, Power BI

Tools: Git, GitHub, FastAPI, Docker, Streamlit, VS Code

Core CS: OOP, DBMS, Operating Systems, Computer Networks

PROJECTS

LitExplorer AI — RAG Research Assistant | *Python, FastAPI, LangChain, FAISS, Ollama* 2025

- Built an end-to-end Retrieval-Augmented Generation pipeline over academic PDFs: pypdf extraction, recursive chunking, `nomic-embed-text` embeddings, FAISS similarity search, and grounded `llama3` generation — running fully on local models with zero external API dependency.
- Engineered token-level streaming end-to-end via FastAPI Server-Sent Events and an async-generator SSE parser on the React client; implemented parallel map-reduce summarization with `asyncio.gather` for structured sectioned output.
- Containerized full stack (Ollama, FastAPI, Nginx + React) with Docker Compose; designed session-scoped FAISS indices for concurrent multi-user querying.
- GitHub: github.com/Akshya-1108/litexplorer

Social-Media Friend Recommendation System | *Python, Scikit-learn, Streamlit, Pandas* Dec 2025

- Developed a content-based recommendation engine that ranks potential friends from structured user profile attributes (age, gender, country, interests) using vectorized similarity scoring.
- Optimized similarity computation by replacing `cosine_similarity` with Scikit-learn's `linear_kernel` on TF-IDF-style sparse vectors; engineered features via one-hot encoding and interest tokenization to improve ranking quality.
- Deployed an interactive Streamlit web app with caching for near real-time inference on repeat queries.
- GitHub: github.com/Akshya-1108/Social-Media-Friend-Recommendation-System

Netflix ETL Pipeline | *Python, Pandas, Snowflake* Mar 2026

- Built a 3-stage ETL pipeline (extract, transform, load) ingesting 8,800+ Netflix records into Snowflake for downstream content-trend analysis; applied null handling, type casting, and duplicate removal for an analysis-ready warehouse table.
- Engineered derived features and automated bulk ingestion via Snowflake's `write_pandas`, significantly reducing load time versus row-by-row inserts.
- GitHub: github.com/Akshya-1108/netflix_etl

EDUCATION

Galgotias University Greater Noida
B.Tech in Computer Science and Engineering CGPA: 7.76 2026

Kendriya Vidyalaya BHEL
Class XII (CBSE) - 75% Class X (CBSE) - 86% 2022 / 2020

ACHIEVEMENTS & CERTIFICATIONS

Solved 100+ algorithmic problems on LeetCode, strengthening problem-solving and data-structure fundamentals.

AWS Academy Machine Learning Certification (2023); Machine Learning Essentials — Udemy (2025).